

DSS5201 DATA VISUALIZATION

WEEK 6

Yuting Huang

NUS DSDS

2024-09-16

Time: Monday Sep 30 7-8pm

Venue: MPSH 2B

Things to bring on the exam day:

- A laptop with the latest Python, VS Code, and Exemplify installed.
- The laptop charger.
- Your NUS matriculation card.

Arrive at least 15 minutes early at the venue for necessary setups (download data sets, check Python packages, etc).

- General information:

<https://nus.atlassian.net/wiki/spaces/DAstudent/pages/22511642/Device+Minimum+System+Requirements>

- Register & join a briefing session on Zoom by Sep 26.

<https://nus.atlassian.net/wiki/spaces/DAstudent/pages/22511675/Common+Briefing+Sessions>

Contact CIT via **citbox25@nus.edu.sg** for any technical issues.

EXAM CONTENT AND FORMAT

- **Content of exam:** All materials covered from Week 1 to Week 6 (inclusive).
- **Format: Open-book, block-internet on Exemplify.**
- You can refer to materials from and beyond our course, but will **not** have access to the internet throughout the exam.
- Make sure you have downloaded data and installed necessary packages before the exam begins.

The total marks available are 30.

- **Part I:** Multiple choice + Fill-in-the-blank questions.
 - Answer questions directly on Exemplify.
 - No submission of Python code is needed.
- **Part II:** Coding questions.
 - Answer questions in a Python Notebook (.ipynb) and render it to HTML (.html).
 - Submit exam on Exemplify, **AND**
 - Submit the Notebook and HTML files to Canvas **immediately after** the exam.

- ① The exam will be available on **Exemplify** from **Sunday, Sep 29 at 8pm**.
 - Only one download is allowed. So ensure you download the exam to the same laptop that you will use during the exam.
- ② Exam data will be available on **Canvas** on **Monday, Sep 30 6:45pm**.
- ③ The following packages are required for the exam:
 - **pandas, numpy, matplotlib, json, pydataset, nycflights13**
 - You may use additional packages, but also need to **ensure** that they are installed properly beforehand.
 - You won't have access to internet once the exam begins.

- ④ The exam begins on **Monday 7pm**, sharp.
- ⑤ During the exam, save your Python Notebook frequently.
- ⑥ The exam **ends at 8pm on Exemplify**.
 - Submit your exam on Exemplify, **AND**
 - Both your python notebook and the rendered HTML to **Canvas** immediately after the exam.
 - On Canvas, the submission window closes at **8:15pm**.

Week 6 & Recess

- Install and update **Python**, **VSCode**, and **Exemplify**
- Verify you've installed and are able load the required packages
`pandas`, `numpy`, `matplotlib`, `json`, `pydataset`, `nycflights13`

Week 7 Sunday 8pm

Download exam from Exemplify

Monday 6:45pm

Download exam data from Canvas

7:00pm
Exam begins

No internet access once exam begins on Exemplify.

Exam begins

8:00pm
Exam ends

Submit exam on Exemplify

Internet access resumes after exam ends on Exemplify.

- Submit `ipynb` and `HTML` to Canvas **immediately after** the exam.
- The Canvas submission folder will close at **8:15pm**

PSET 1: PRACTICE QUESTIONS

PSet 1 will be available on Canvas.

- Practice coding questions on data import, data manipulation, and data joins.
- Use Python notebook to answer questions and render it into HTML at the end.
- Due on **this Friday, Sep 20 by 11:59pm**.

Best to complete them independently, in an environment similar to the exam setting.

Data transformation

Week 4

- `query()`, `sort_values()`, `rename()`, `groupby()`, ...

Tidy data (data reshaping)

Week 5

- `melt()`, `stack()`, `pivot()`, and `pivot_table()`.

Relational data

Week 6

Data never arrive in the condition that we need them.

They need to be reshaped and reformatted.

“Tidy” Table

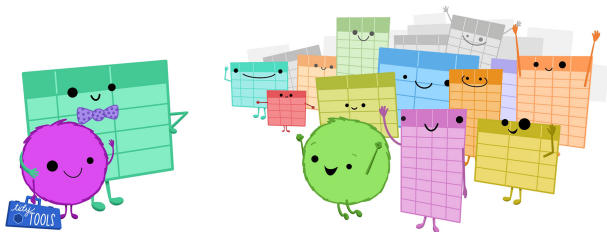
| Business Unit | Year | Quarter | Budget |
|---------------|------|---------|-----------|
| Sales | 2000 | Q1 | 2,500,000 |
| Marketing | 2000 | Q1 | 1,000,000 |
| Sales | 2000 | Q2 | 2,750,000 |
| Marketing | 2000 | Q2 | 1,250,000 |
| Sales | 2000 | Q3 | 3,000,000 |
| Marketing | 2000 | Q3 | 4,000,000 |
| Sales | 2000 | Q4 | 2,000,000 |
| Marketing | 2000 | Q4 | 500,000 |
| Sales | 2001 | Q1 | 2,500,000 |
| Marketing | 2001 | Q1 | 1,500,000 |

“UnTidy” Table

| Past and projected budgets for WidgetCo.'s Sales and Marketing Org. | | | | | | |
|---|-----------|-----------|-----------|-----------|-----------|-----------|
| Contact JDoe@widgets.ca for more information. | | | | | | |
| | Year | | | | | |
| Business Unit | 2000 | | | | 2001 | |
| | Q1 | Q2 | Q3 | Q4 | Q1 | Q2* |
| Sales | 2,500,000 | 2,750,000 | 3,000,000 | 2,000,000 | 2,500,000 | 3,000,000 |
| Marketing | 1,000,000 | 1,250,000 | 4,000,000 | 500,000 | 1,500,000 | 1,750,000 |
| double check this number - JD | | | | | | |
| *Projected Numbers | | | | | | |

- The tidy table is ready for use. The untidy table is not.

WHEN ONE TABLE IS NOT ENOUGH



When working with real-world data, you will often find that data are stored across **multiple** files or data frames.

- Typically, these tables have to be combined to answer the questions we are interested in.
- Many tables of data are called **relational data**.

RELATIONAL DATA

WHEN ONE TABLE IS NOT ENOUGH

| restaurant | | | | health inspections | | | | | rating | | |
|--------------|--------|----------------|------------|--------------------|--------|-----------------|-----------|-------|--------------|--------|-------|
| name | id | address | type | name | id | inspection_date | inspector | score | name | id | stars |
| Taco Stand | AH13JK | 1 Main St. | Mexican | Taco Stand | AH13JK | 2018-08-21 | Sheila | 97 | Taco Stand | AH13JK | 4.9 |
| Pho Place | JJ29JJ | 192 Street Rd. | Vietnamese | Pho Place | JJ29JJ | 2018-03-12 | D'eonte | 98 | Pho Place | JJ29JJ | 4.8 |
| Taco Stand | XJ11AS | 18 W. East St. | Fusion | Pho Place | JJ29JJ | 2018-01-02 | Monica | 66 | Taco Stand | XJ11AS | 4.2 |
| Pizza Heaven | CI21AA | 711 K Ave. | Italian | Taco Stand | XJ11AS | 2018-12-16 | Mark | 43 | Pizza Heaven | CI21AA | 4.7 |
| | | | | Pizza Heaven | CI21AA | 2018-08-21 | Anh | 99 | | | |

Consider a town with a number of restaurants.

Across multiple data files, we have information on

- Location and type of cuisine.
- Health and safety inspections results.
- Online ratings on the restaurant.

ADVANTAGES OF RELATIONAL DATA

Storing data across multiple files has a number of benefits:

- **Efficient data storage:** Limit the need to repeat information.
- **Easier data updates:** If we need to update information, we can make the change in a single file.
- **Privacy:** We can restrict access to some of the data to ensure only those who should have access are able to read the data.

Today, we work with **five** related tables from about flights in the New York City in 2013.

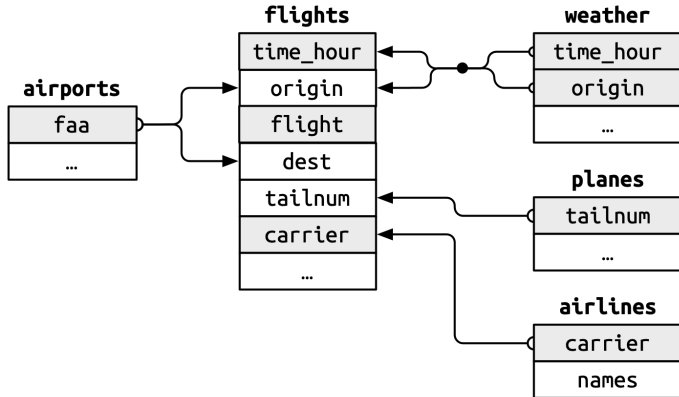
- ① flights: All flights that departed New York City in 2013.
- ② airlines: Carrier name and its abbreviated code.
- ③ airports: Information about airports.
- ④ planes: Plane's tailnum found in the FAA aircraft registry.
- ⑤ weather: Weather at each airport in New York at each hour.

The data are available in the [nycflights13](#) package.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

import nycflights13
from nycflights13 import flights
from nycflights13 import airlines
from nycflights13 import airports
from nycflights13 import planes
from nycflights13 import weather
```

Here is a diagram (database schema) that identifies the connections between tables:



In airlines, we can look up the full carrier name from its abbreviated code.

```
airlines.head()
```

| | carrier | name |
|---|---------|------------------------|
| 0 | 9E | Endeavor Air Inc. |
| 1 | AA | American Airlines Inc. |
| 2 | AS | Alaska Airlines Inc. |
| 3 | B6 | JetBlue Airways |
| 4 | DL | Delta Air Lines Inc. |

airports gives information about each airport, with a unique faa airport code.

```
airports.head()
```

| | faa | name | lat | ... | tz | dst | |
|---|-----|-------------------------------|-----------|-----|----|-----|-------------|
| 0 | 04G | Lansdowne Airport | 41.130472 | ... | -5 | A | America/New |
| 1 | 06A | Moton Field Municipal Airport | 32.460572 | ... | -6 | A | America/CL |
| 2 | 06C | Schaumburg Regional | 41.989341 | ... | -6 | A | America/CL |
| 3 | 06N | Randall Airport | 41.431912 | ... | -5 | A | America/New |
| 4 | 09J | Jekyll Island Airport | 31.074472 | ... | -5 | A | America/New |

```
[5 rows x 8 columns]
```

planes provides information about each plane, with a unique tailnum (tail number).

```
planes.head()
```

| | tailnum | year | type | ... | seats | speed | engine |
|---|---------|--------|-------------------------|-----|-------|-------|-----------|
| 0 | N10156 | 2004.0 | Fixed wing multi engine | ... | 55 | NaN | Turbo-fan |
| 1 | N102UW | 1998.0 | Fixed wing multi engine | ... | 182 | NaN | Turbo-fan |
| 2 | N103US | 1999.0 | Fixed wing multi engine | ... | 182 | NaN | Turbo-fan |
| 3 | N104UW | 1999.0 | Fixed wing multi engine | ... | 182 | NaN | Turbo-fan |
| 4 | N10575 | 2002.0 | Fixed wing multi engine | ... | 55 | NaN | Turbo-fan |

```
[5 rows x 9 columns]
```

weather provides hourly meterological data for the three airports in New York.

```
weather.head()
```

| | origin | year | month | day | ... | precip | pressure | visib | time_l |
|---|--------|------|-------|-----|-----|--------|----------|-------|------------------|
| 0 | EWR | 2013 | 1 | 1 | ... | 0.0 | 1012.0 | 10.0 | 2013-01-01T06:00 |
| 1 | EWR | 2013 | 1 | 1 | ... | 0.0 | 1012.3 | 10.0 | 2013-01-01T07:00 |
| 2 | EWR | 2013 | 1 | 1 | ... | 0.0 | 1012.5 | 10.0 | 2013-01-01T08:00 |
| 3 | EWR | 2013 | 1 | 1 | ... | 0.0 | 1012.2 | 10.0 | 2013-01-01T09:00 |
| 4 | EWR | 2013 | 1 | 1 | ... | 0.0 | 1011.9 | 10.0 | 2013-01-01T10:00 |

```
[5 rows x 15 columns]
```

The variable that connects a pair of data sets are called **keys**.

- A variable (or a *minimal* set of variables) that uniquely identifies an observation in a data frame.

In the database schema,

- In the `planes` table, `tailnum` is the key variable.
- In the `weather` table, each observation is uniquely identified by a set of two variables: `time_hour`, and `origin`.

Each data join involves a pair of keys: Primary key and foreign key.

- **Primary key** uniquely identifies an observation in the same table.
- For example, carrier is the primary key for airlines:

```
airlines.head()
```

| | carrier | name |
|---|---------|------------------------|
| 0 | 9E | Endeavor Air Inc. |
| 1 | AA | American Airlines Inc. |
| 2 | AS | Alaska Airlines Inc. |
| 3 | B6 | JetBlue Airways |
| 4 | DL | Delta Air Lines Inc. |

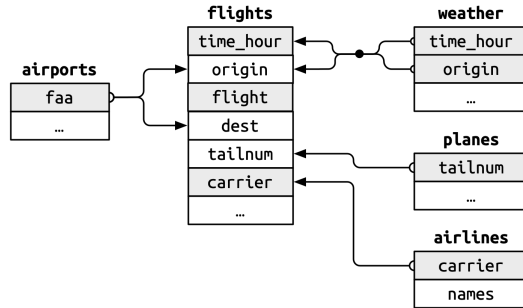
When more than one variable is needed, the key is called a **compound key**.

- origin and time_hour are the compound key for the weather table:

```
weather.head()
```

| | origin | year | month | day | ... | precip | pressure | visib | time_l |
|---|--------|------|-------|-----|-----|--------|----------|-------|------------------|
| 0 | EWR | 2013 | 1 | 1 | ... | 0.0 | 1012.0 | 10.0 | 2013-01-01T06:00 |
| 1 | EWR | 2013 | 1 | 1 | ... | 0.0 | 1012.3 | 10.0 | 2013-01-01T07:00 |
| 2 | EWR | 2013 | 1 | 1 | ... | 0.0 | 1012.5 | 10.0 | 2013-01-01T08:00 |
| 3 | EWR | 2013 | 1 | 1 | ... | 0.0 | 1012.2 | 10.0 | 2013-01-01T09:00 |
| 4 | EWR | 2013 | 1 | 1 | ... | 0.0 | 1011.9 | 10.0 | 2013-01-01T10:00 |

```
[5 rows x 15 columns]
```

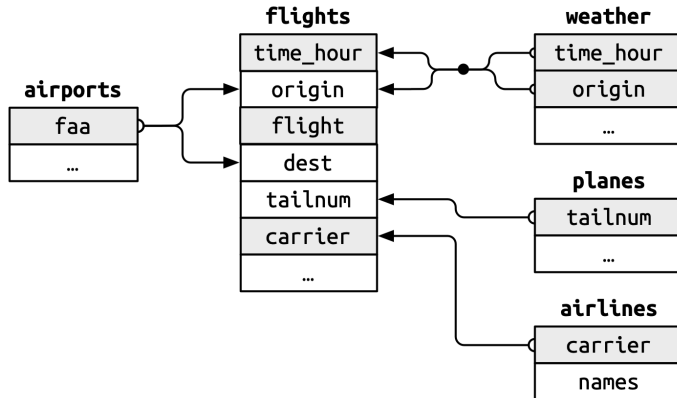


Foreign key is the counterpart of primary key. It uniquely identifies an observation in a different table.

- `flights["carrier"]` is a foreign key that corresponds to the primary key `airlines["carrier"]`.
- A variable can be a primary and a foreign key at the same time.

PRIMARY AND FOREIGN KEYS

These relationship can be summarized visually in the database schema.



Once we've identified primary key(s), it is a good practice to verify that they can indeed **uniquely** identify each observation.

- One way to do that is to use `value_counts()` on the key and look for entries where the count is greater than 1.

```
counts = airlines["carrier"].value_counts().reset_index(name = "n")
counts[counts["n"] > 1]
```

Empty DataFrame

Columns: [carrier, n]

Index: []

- That means, carrier can uniquely identify observations in airlines.

Let's turn to check the compound key for the weather table.

- From the schema, it appears that `time_hour` and `origin` can identify an observation in weather.
- We can check if that's true.

```
counts = weather[["time_hour", "origin"]].value_counts()  
counts = counts.reset_index(name = "n")  
counts[counts["n"] > 1]
```

Empty DataFrame

Columns: [time_hour, origin, n]

Index: []

We should also check for missing values in primary keys.

- If the key is missing, the variable cannot identify any observation.

```
checks = planes[planes["tailnum"].isna()]
checks.head()
```

Empty DataFrame

Columns: [tailnum, year, type, manufacturer, model, engines, seats, speed,
Index: []

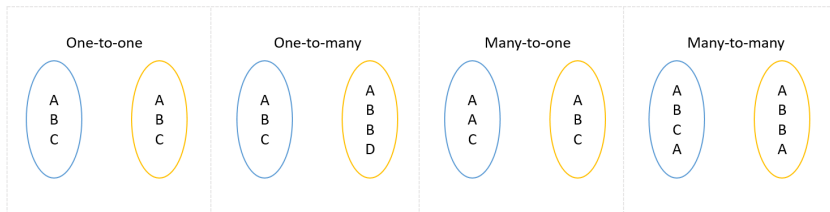
```
checks = weather[weather["time_hour"].isna() | weather["origin"].isna()]
checks.head()
```

Empty DataFrame

Columns: [origin, year, month, day, hour, temp, dewp, humid, wind_dir, wind

Index: []

- We should also check these for other tables in the database!

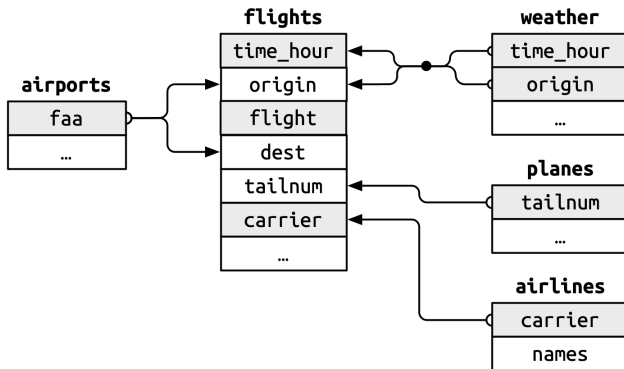


A primary key and the corresponding foreign key forms a **relation**.

- Ideally, relationships are **one-to-one**.
- In real-life data sets, relations are typically **one-to-many** or **many-to-one**:
 - E.g., each flight has one plane, but each plane flies many flights.
- Relations can also be **many-to-many**:
 - Each airline flies to many airports, each airport hosts many airlines.

To work with relational data, we need functions that works with a pair of tables.

- **Mutating joins:** Add new variables to one data frame from matching observations in another data frame.
- **Filtering joins:** Filter observations from one data frame based on whether they can be matched to an observation in another data frame.
- **Inequality joins.**



Let's combine a pair of tables using **mutating join**.

- `flights` and `airlines` via `carrier`.

To ease demonstration, let's first create a narrower data frame that contains fewer variables.

- We name it as `flights2`.

```
flights2 = flights[["time_hour", "origin", "dest", "tailnum", "carrier"]]  
flights2.head()
```

| | time_hour | origin | dest | tailnum | carrier |
|---|----------------------|--------|------|---------|---------|
| 0 | 2013-01-01T10:00:00Z | EWB | IAH | N14228 | UA |
| 1 | 2013-01-01T10:00:00Z | LGA | IAH | N24211 | UA |
| 2 | 2013-01-01T10:00:00Z | JFK | MIA | N619AA | AA |
| 3 | 2013-01-01T10:00:00Z | JFK | BQN | N804JB | B6 |
| 4 | 2013-01-01T11:00:00Z | LGA | ATL | N668DN | DL |

Let's also take a look at the airlines table.

```
airlines.head()
```

| | carrier | name |
|---|---------|------------------------|
| 0 | 9E | Endeavor Air Inc. |
| 1 | AA | American Airlines Inc. |
| 2 | AS | Alaska Airlines Inc. |
| 3 | B6 | JetBlue Airways |
| 4 | DL | Delta Air Lines Inc. |

- **carrier** is a primary key in airlines.
- We can join the airlines and flights2 tables via **carrier**.

Notice that we use a left join (`how = "left"`) on `carrier`.

- The name of the airline is added to the right of the `flights2` table.

```
df = flights2.merge(airlines, how = "left", on = "carrier")
df.head()
```

| | time_hour | origin | dest | tailnum | carrier | name |
|---|----------------------|--------|------|---------|---------|-----------------------|
| 0 | 2013-01-01T10:00:00Z | EWR | IAH | N14228 | UA | United Air Lines Inc |
| 1 | 2013-01-01T10:00:00Z | LGA | IAH | N24211 | UA | United Air Lines Inc |
| 2 | 2013-01-01T10:00:00Z | JFK | MIA | N619AA | AA | American Airlines Inc |
| 3 | 2013-01-01T10:00:00Z | JFK | BQN | N804JB | B6 | JetBlue Airways |
| 4 | 2013-01-01T11:00:00Z | LGA | ATL | N668DN | DL | Delta Air Lines Inc |

In the following, we will learn four types of **mutating joins**.

- All of them can be performed using the `merge()` function. We just need to specify the `how` argument inside of it.
 - `how = "left"`: Left join
 - `how = "right"`: Right join
 - `how = "inner"`: Inner join
 - `how = "outer"`: Outer join

To understand how they work, let's create simpler data sets and use visual representations.

| x | | y | |
|-----|-------|-----|-------|
| key | var_x | key | var_y |
| 1 | x1 | 1 | y1 |
| 2 | x2 | 2 | y2 |
| 3 | x3 | 4 | y3 |

```
# Create data frames x and y
x = pd.DataFrame({
    "key": [1, 2, 3],
    "val_x": ["x1", "x2", "x3"]
})

y = pd.DataFrame({
    "key": [1, 2, 4],
    "val_y": ["y1", "y2", "y3"]
})
```

| x | | y | |
|-----|-------|-----|-------|
| key | var_x | key | var_y |
| 1 | x1 | 1 | y1 |
| 2 | x2 | 2 | y2 |
| 3 | x3 | 4 | y3 |

- The colored column represents the **key** variable.
- The grey column represents the **value**.

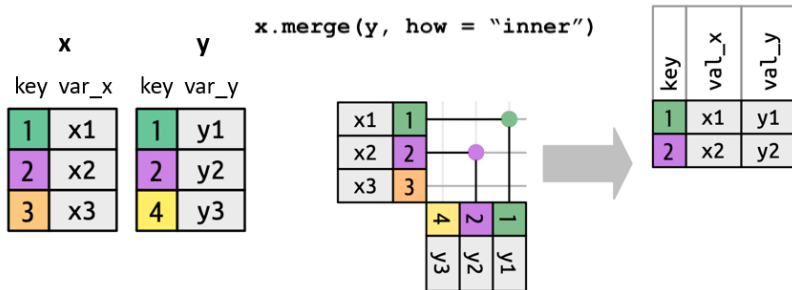
For simplicity, we only show the case with one key variable. But the idea generalizes to multiple keys and multiple values.

A join is a way of connecting each row in table x to zero, one, or more rows in table y.

| | | | | |
|----|---|----|----|----|
| x1 | 1 | | | |
| x2 | 2 | | | |
| x3 | 3 | | | |
| | | 4 | 2 | 1 |
| | | y3 | y2 | y1 |

- If you look closely, you may notice that we switched the order of the key and value columns in table x.
- This is to emphasize that joins matches based on the **key** variable.

In an actual join, matches will be indicated with dots.



- Number of dots = number of matches.
- Different types of joins will result in different number of rows.

Inner join is the simplest type of data joins.

- Matches pairs of observations whenever their keys are equal.
- Keeps observations that appear in **both** tables, and remove the unmatched ones.

```
# Inner join  
x.merge(y, how = "inner", on = "key")
```

| | key | val_x | val_y |
|---|-----|-------|-------|
| 0 | 1 | x1 | y1 |
| 1 | 2 | x2 | y2 |

LEFT, RIGHT, AND OUTER JOINS

These join types keep observations that appear in **at least one** of the two tables.

- Left join: Keeps all rows in x , including those unmatched to y .
- Right join: Keeps all rows in y , including those unmatched to x .
- Outer join: Keeps all rows in both tables, regardless of matches.
- Cross join: Creates the cartesian product from both x and y .

These joins work by adding “virtual” observations to each table. The matched observations have their original values, the unmatched ones are filled with NaN.

```
# Left join
x.merge(y, how = "left", on = "key")
```

| | key | val_x | val_y |
|---|-----|-------|-------|
| 0 | 1 | x1 | y1 |
| 1 | 2 | x2 | y2 |
| 2 | 3 | x3 | NaN |

```
# Right join
x.merge(y, how = "right", on = "key")
```

| | key | val_x | val_y |
|---|-----|-------|-------|
| 0 | 1 | x1 | y1 |
| 1 | 2 | x2 | y2 |
| 2 | 4 | NaN | y3 |

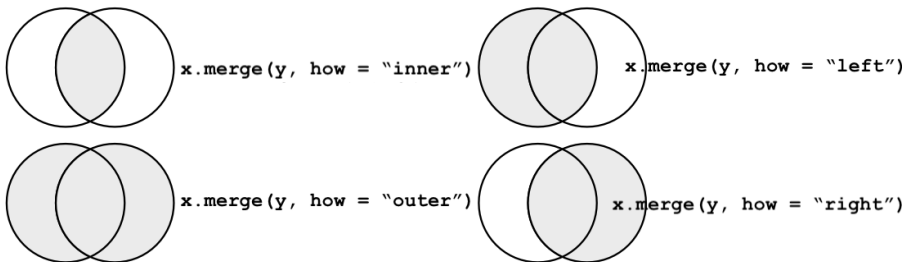
```
# Outer join
x.merge(y, how = "outer", on = "key")
```

| | key | val_x | val_y |
|---|-----|-------|-------|
| 0 | 1 | x1 | y1 |
| 1 | 2 | x2 | y2 |
| 2 | 3 | x3 | NaN |
| 3 | 4 | NaN | y3 |

```
# Cross join (much less often used)
x.merge(y, how = "cross")
```

| | key_x | val_x | key_y | val_y |
|---|-------|-------|-------|-------|
| 0 | 1 | x1 | 1 | y1 |
| 1 | 1 | x1 | 2 | y2 |
| 2 | 1 | x1 | 4 | y3 |
| 3 | 2 | x2 | 1 | y1 |
| 4 | 2 | x2 | 2 | y2 |
| 5 | 2 | x2 | 4 | y3 |

USE LEFT JOIN AS YOUR DEFAULT JOIN



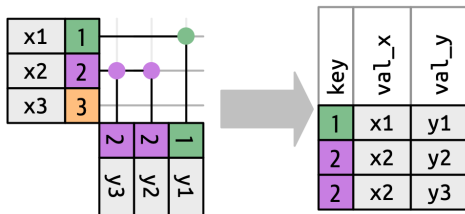
The most common join is **left join**, as it preserves the original observation even when there isn't a match.

- **Left join should be your default join**, unless you have a strong reason to prefer one of the others.

So far, we've explored what happens if a row in x matches zero or one row in y .

This is not always the case.

- 1 If one table has duplicated keys, then the matching row will be duplicated as well.

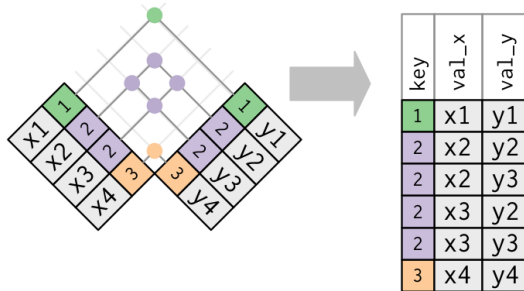



```
x = pd.DataFrame({
    "key": [1, 2, 3],
    "val_x": ["x1", "x2", "x3"]
})
y = pd.DataFrame({
    "key": [1, 2, 2],
    "val_y": ["y1", "y2", "y3"]
})
x.merge(y, how = "inner", on = "key")
```

| | key | val_x | val_y |
|---|-----|-------|-------|
| 0 | 1 | x1 | y1 |
| 1 | 2 | x2 | y2 |
| 2 | 2 | x2 | y3 |

② If both table have duplicated keys, you get all possible combinations, the Cartesian product:

- However, this is usually a data error.
- In most cases, need to have **unique keys** for at least one of your tables.



```
x = pd.DataFrame({
    "key": [1, 2, 2, 3],
    "val_x": ["x1", "x2", "x3", "x4"]
})
y = pd.DataFrame({
    "key": [1, 2, 2, 3],
    "val_y": ["y1", "y2", "y3", "y4"]
})
x.merge(y, how = "left", on = "key")
```

| | key | val_x | val_y |
|---|-----|-------|-------|
| 0 | 1 | x1 | y1 |
| 1 | 2 | x2 | y2 |
| 2 | 2 | x2 | y3 |
| 3 | 2 | x3 | y2 |
| 4 | 2 | x3 | y3 |
| 5 | 3 | x4 | y4 |

Many-to-many joins are particularly **problematic**

- Because they can result in a **size explosion** of the object returned from the join.
- This will have a large impact on the performance of your code.

BACK TO THE NYC FLIGHTS DATA

Let's return to the flights data, flights2.

```
flights2.head(10)
```

| | time_hour | origin | dest | tailnum | carrier |
|---|----------------------|--------|------|---------|---------|
| 0 | 2013-01-01T10:00:00Z | EWR | IAH | N14228 | UA |
| 1 | 2013-01-01T10:00:00Z | LGA | IAH | N24211 | UA |
| 2 | 2013-01-01T10:00:00Z | JFK | MIA | N619AA | AA |
| 3 | 2013-01-01T10:00:00Z | JFK | BQN | N804JB | B6 |
| 4 | 2013-01-01T11:00:00Z | LGA | ATL | N668DN | DL |
| 5 | 2013-01-01T10:00:00Z | EWR | ORD | N39463 | UA |
| 6 | 2013-01-01T11:00:00Z | EWR | FLL | N516JB | B6 |
| 7 | 2013-01-01T11:00:00Z | LGA | IAD | N829AS | EV |
| 8 | 2013-01-01T11:00:00Z | JFK | MCO | N593JB | B6 |
| 9 | 2013-01-01T11:00:00Z | LGA | ORD | N3ALAA | AA |

There are several ways to specify the key variables.

- 1 Specify the argument `on = "key_variable"`.

```
df = flights2.merge(airlines, how = "left", on = "carrier")
df.head()
```

| | time_hour | origin | dest | tailnum | carrier | name |
|---|----------------------|--------|------|---------|---------|-----------------------|
| 0 | 2013-01-01T10:00:00Z | EWR | IAH | N14228 | UA | United Air Lines Inc |
| 1 | 2013-01-01T10:00:00Z | LGA | IAH | N24211 | UA | United Air Lines Inc |
| 2 | 2013-01-01T10:00:00Z | JFK | MIA | N619AA | AA | American Airlines Inc |
| 3 | 2013-01-01T10:00:00Z | JFK | BQN | N804JB | B6 | JetBlue Airways |
| 4 | 2013-01-01T11:00:00Z | LGA | ATL | N668DN | DL | Delta Air Lines Inc |

- ② Leave the `on` argument empty. Then the `merge()` function will use the common variable(s) in the two tables.

- In the example below, the two tables are joined via `carrier`.

```
df = flights2.merge(airlines, how = "left")
df.head()
```

| | time_hour | origin | dest | tailnum | carrier | name |
|---|----------------------|--------|------|---------|---------|-----------------------|
| 0 | 2013-01-01T10:00:00Z | EWR | IAH | N14228 | UA | United Air Lines Inc |
| 1 | 2013-01-01T10:00:00Z | LGA | IAH | N24211 | UA | United Air Lines Inc |
| 2 | 2013-01-01T10:00:00Z | JFK | MIA | N619AA | AA | American Airlines Inc |
| 3 | 2013-01-01T10:00:00Z | JFK | BQN | N804JB | B6 | JetBlue Airways |
| 4 | 2013-01-01T11:00:00Z | LGA | ATL | N668DN | DL | Delta Air Lines Inc |

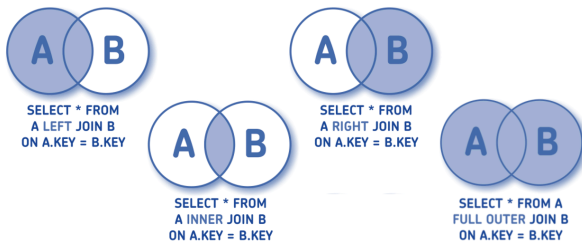
- ③ If the names of the key variables are different in two tables, specify the `left_on` and `right_on` parameters.

```
df = flights2.merge(
    airports, how = "left", left_on = "dest", right_on = "faa"
)
df.head()
```

| | time_hour | origin | dest | tailnum | ... | alt | tz | dst | |
|---|----------------------|--------|------|---------|-----|--------|------|-----|---------|
| 0 | 2013-01-01T10:00:00Z | EWR | IAH | N14228 | ... | 97.0 | -6.0 | A | America |
| 1 | 2013-01-01T10:00:00Z | LGA | IAH | N24211 | ... | 97.0 | -6.0 | A | America |
| 2 | 2013-01-01T10:00:00Z | JFK | MIA | N619AA | ... | 8.0 | -5.0 | A | America |
| 3 | 2013-01-01T10:00:00Z | JFK | BQN | N804JB | ... | NaN | NaN | NaN | |
| 4 | 2013-01-01T11:00:00Z | LGA | ATL | N668DN | ... | 1026.0 | -5.0 | A | America |

[5 rows x 13 columns]

SQL JOINS



The [pandas user guide](#) provides the full documentation for `merge()`.

The translation to SQL is straightforward:

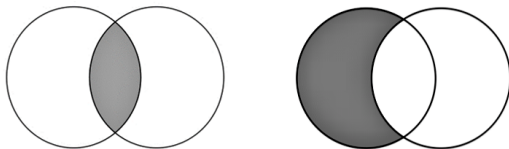
| Pandas | SQL |
|--|---|
| <code>x.merge(y, on = "key", how = "left")</code> | <code>SELECT * FROM x LEFT JOIN y ON (key)</code> |
| <code>x.merge(y, on = "key", how = "right")</code> | <code>SELECT * FROM x RIGHT JOIN y ON (key)</code> |
| <code>x.merge(y, on = "key", how = "inner")</code> | <code>SELECT * FROM x INNER JOIN y ON (key)</code> |
| <code>x.merge(y, on = "key", how = "outer")</code> | <code>SELECT * FROM x FULL OUTER JOIN y ON (key)</code> |

FILTERING JOINS

Filtering joins match observations in the same way as mutating joins, but affect the observations in the final table.

There are two types:

- `isin()`: Keep all observation in `x` that has a match in `y`.
- Negation of `isin()`: Remove all observation in `x` that has a match in `y`.

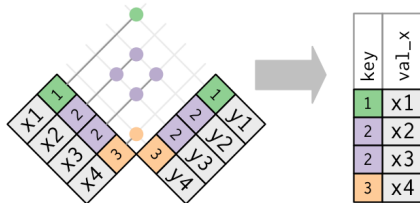


KEEPING THE MATCHED RECORDS

`isin()` keeps only the **matched** observations in `x`.



If there are duplicated keys in `x`, all those rows will be kept.



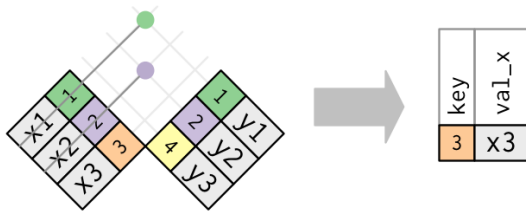
- Find all flights that flew to the most popular destinations:

```
# The most popular destination airports
top = flights["dest"].value_counts().nlargest(1).reset_index(name = "n")
# Keeping flights that flew to that destination
flights2[flights2["dest"].isin(top["dest"])]
```

| | time_hour | origin | dest | tailnum | carrier |
|--------|----------------------|--------|------|---------|---------|
| 5 | 2013-01-01T10:00:00Z | EWR | ORD | N39463 | UA |
| 9 | 2013-01-01T11:00:00Z | LGA | ORD | N3ALAA | AA |
| 25 | 2013-01-01T11:00:00Z | EWR | ORD | N9EAMQ | MQ |
| 38 | 2013-01-01T11:00:00Z | LGA | ORD | N3CYAA | AA |
| 57 | 2013-01-01T12:00:00Z | LGA | ORD | N4WNAA | AA |
| ... | ... | ... | ... | ... | ... |
| 336645 | 2013-09-30T23:00:00Z | LGA | ORD | N4XBAA | AA |
| 336669 | 2013-10-01T00:00:00Z | LGA | ORD | N853UA | UA |
| 336675 | 2013-10-01T00:00:00Z | EWR | ORD | N511MQ | MQ |
| 336696 | 2013-10-01T00:00:00Z | JFK | ORD | N298JB | B6 |
| 336709 | 2013-10-01T00:00:00Z | LGA | ORD | N434AA | AA |

REMOVING THE MATCHED RECORDS

The function can also be used to keep only the **unmatched** records.



- It is useful for diagnosing join mismatches.

- Identify the flights that do not have a match in planes:

```
# Find the unmatched records
unmatched = flights[~flights["tailnum"].isin(planes["tailnum"])]
# Count each unmatched tailnum
counts = unmatched["tailnum"].value_counts().reset_index(name = "n")
counts.head()
```

| | tailnum | n |
|---|---------|-----|
| 0 | N725MQ | 575 |
| 1 | N722MQ | 513 |
| 2 | N723MQ | 507 |
| 3 | N713MQ | 483 |
| 4 | N735MQ | 396 |

The data we have seen in class have been cleaned up so you have as few problems as possible.

Your own data is unlikely to be so nice.

So there are a few things you should do with your own data to make your joins go more smoothly.

- 1 Identify the primary keys in each variable.
- 2 Check that none of the variables in the primary key are missing. If a value is missing, it cannot identify an observation.
- 3 Check that foreign keys match primary keys in another table.

So far, we've also only seen joins where rows are matched if the x key equals the y key.

Now we will relax this restriction.

- We shall introduce **inequality joins**.
- It matches rows based on an **inequality condition** between the keys.
- It is supported in SQL. Let's see how it can be done in pandas.

```
# Create two data frames
sales = pd.DataFrame({
    "sales_date": pd.to_datetime(["2024-09-01", "2024-09-03",
                                   "2024-09-14", "2024-09-17"])
})

promos = pd.DataFrame({
    "promo_date": pd.to_datetime(["2024-09-09", "2024-09-15"]),
    "promo_price": [179, 179]
})
```

TABLES ON SALES AND PROMO

sales

| | sales_date |
|---|------------|
| 0 | 2024-09-01 |
| 1 | 2024-09-03 |
| 2 | 2024-09-14 |
| 3 | 2024-09-17 |

promos

| | promo_date | promo_price |
|---|------------|-------------|
| 0 | 2024-09-09 | 179 |
| 1 | 2024-09-15 | 179 |

- We've learned about **inner join** – matching rows where sales_date equals promo_date.

```
sales.merge(promos, how = "inner",  
            left_on = "sales_date", right_on = "promo_date")
```

Empty DataFrame

Columns: [sales_date, promo_date, promo_price]

Index: []

- There is no match since sales_date and promo_date does not exactly equal.

- **Inequality join:** Matching rows where sales_date occurs after promo_date.

```
df = sales.merge(promos, how = "cross")  
df.query("sales_date >= promo_date")
```

| | sales_date | promo_date | promo_price |
|---|------------|------------|-------------|
| 4 | 2024-09-14 | 2024-09-09 | 179 |
| 6 | 2024-09-17 | 2024-09-09 | 179 |
| 7 | 2024-09-17 | 2024-09-15 | 179 |

A two-step operation:

- Uses a **cross join**, which joins all rows from both tables.
- Filter the resulting data frame to keep only the rows where sales_date occurs after promo_date.

ADDITIONAL PACKAGE: PYJANITOR

A more elegant (and efficient) way exists, but we need the `pyjanitor` package.

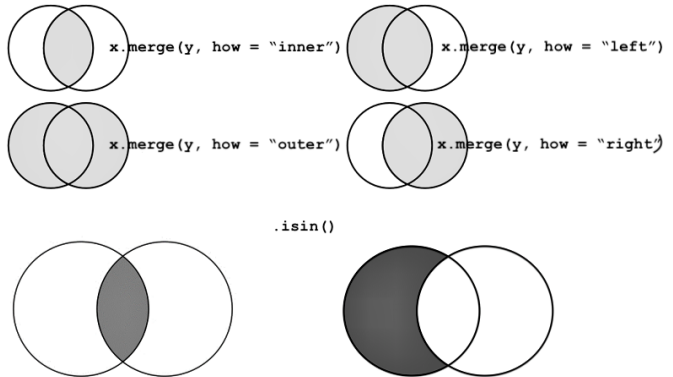
- The function we will be using is `conditional_join()`.

```
import janitor as jn
sales.conditional_join(promos, ("sales_date", "promo_date", ">="),
                        how = "inner")
```

| | sales_date | promo_date | promo_price |
|---|------------|------------|-------------|
| 0 | 2024-09-14 | 2024-09-09 | 179 |
| 1 | 2024-09-17 | 2024-09-09 | 179 |
| 2 | 2024-09-17 | 2024-09-15 | 179 |

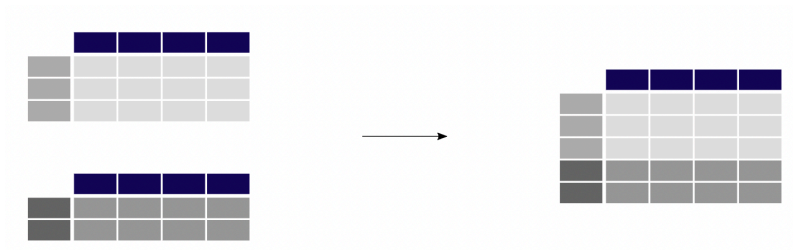
SUMMARY OF RELATIONAL DATA

- **Mutating joins:** Match by key variables and keep columns of both inputs.
- **Filtering joins:** Match by key variables and keep columns of the first input.
- **Inequality joins:** Relax the restrictions on keys being equal.



CONCATINATION

If we have two or more data frames with the same index or the same columns, we can **concatenate** them using `pd.concat()`.



- Create two data frames with the same columns:

```
# Create two data frames
sales = pd.DataFrame({
    'sales_date': pd.to_datetime(["2024-09-01", "2024-09-03",
                                   "2024-09-14", "2024-09-17"])
})

sales1 = pd.DataFrame({
    'sales_date': pd.to_datetime(["2024-08-01"])
})
```

① Concatenate data frames **vertically** (i.e., appending data frames).

- `axis = 0` specifies that the tables should be concatenated along rows.
- `ignore_index = True` tells the function to ignore the original row indices, and create a new set of index starting from 0.

```
## Concatenate the data frames vertically  
pd.concat([sales, sales1], axis = 0, ignore_index = True)
```

```
   sales_date  
0 2024-09-01  
1 2024-09-03  
2 2024-09-14  
3 2024-09-17  
4 2024-08-01
```

② Concatenate data frames **horizontally**.

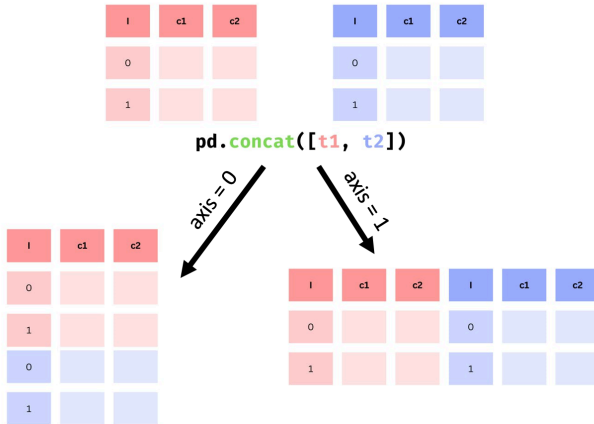
- `axis = 1` specifies that the tables should be concatenated along columns.

```
prices = pd.DataFrame({  
    "quantity": [3, 9, 1, 2],  
    "price": [209, 209, 179, 179]  
})
```

```
## Concatenate the data frames horizontally  
pd.concat([sales, prices], axis = 1)
```

| | sales_date | quantity | price |
|---|------------|----------|-------|
| 0 | 2024-09-01 | 3 | 209 |
| 1 | 2024-09-03 | 9 | 209 |
| 2 | 2024-09-14 | 1 | 179 |
| 3 | 2024-09-17 | 2 | 179 |

SUMMARY OF CONCATINATION



A REVIEW

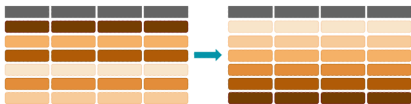


Customarily, we import them as:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

- `pandas` functions can handle many different types of files.
 - `read_csv()`, `read_excel()`, `read_html()`, ...
- For json files, we use the `json` package.
- For API requests, we use the `requests` package.
- Data can be stored in packages too. So far we've used `pydataset` and `nycflights13` for that.

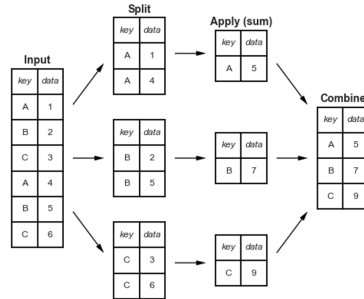
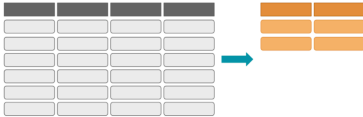
Row manipulations



Column manipulations



Groups and summaries

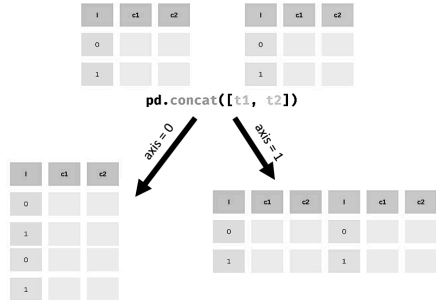
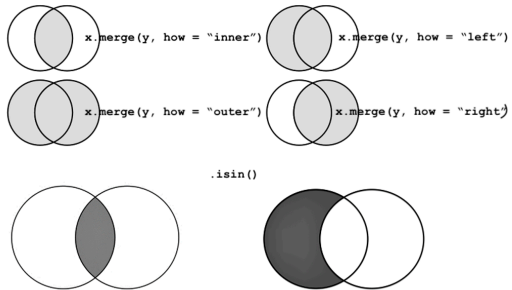


Wide to long: `melt()`

| country | year | cases |
|-------------|------|--------|
| Afghanistan | 1999 | 745 |
| Afghanistan | 2000 | 2666 |
| Brazil | 1999 | 37737 |
| Brazil | 2000 | 80488 |
| China | 1999 | 212258 |
| China | 2000 | 213766 |

Long to wide: `pivot()`

| country | year | type | count |
|-------------|------|------------|------------|
| Afghanistan | 1999 | cases | 745 |
| Afghanistan | 1999 | population | 19987071 |
| Afghanistan | 2000 | cases | 2666 |
| Afghanistan | 2000 | population | 20595360 |
| Brazil | 1999 | cases | 37737 |
| Brazil | 1999 | population | 172006362 |
| Brazil | 2000 | cases | 80488 |
| Brazil | 2000 | population | 174504898 |
| China | 1999 | cases | 212258 |
| China | 1999 | population | 1272915272 |
| China | 2000 | cases | 213766 |
| China | 2000 | population | 1280428583 |



- For specific tasks, functions in packages like **pyjanitor** can be useful as well.

DATA VISUALIZATION (BASICS)

So far, we've primarily used `pandas` and `matplotlib` for visualization.

- Histogram: For quantitative (continuous) variables.
- Bar chart: For qualitative (categorical) variables.
- Line chart: Time-series data.

To visualize more more than two variables on a chart, it's easier with `seaborn`, especially when we have **tidy data**.

- We will cover more about visualization after the midterm exam.